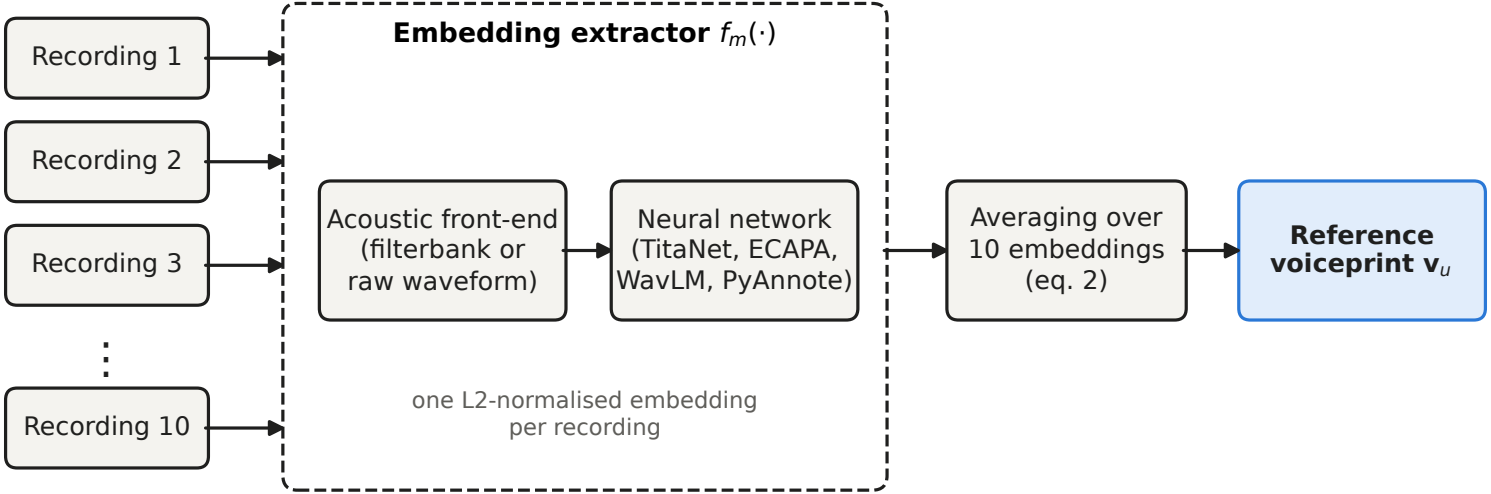
Recording 1
Recording 2
Recording 3
Recording 10
Embedding extractor $f_m(\cdot)$
Acoustic front-end (filterbank or raw waveform)
Neural network (TitaNet, ECAPA, WavLM, PyAnnote)
one L2-normalised embedding per recording
Averaging over 10 embeddings (eq. 2)
Reference voiceprint $\mathbf{v}_u$